


Sure to Sorry

Safety amplification

SafeInfer (inference time intervention)
Safety guided decoding AAAI 2028





Function vectors $\rightarrow$ safety amplification

Model arithmetic $\rightarrow$ safety guided decoding

Function vectors

Function vectors - causal mediation analysis
(Todd et al., ICLR 2024)

- Transformer model f

- ICL prompt $p_i^t \in P_t$ dataset representing task t

$$p_i^t = [(x_{i1}, y_{i1}), \dots, (x_{iN}, y_{iN}), x_{iq}] \rightarrow y_{iq}$$

$$\tilde{p}_i^t = [(x_{i1}, \tilde{y}_{i1}), \dots, (x_{iN}, \tilde{y}_{iN}), x_{iq}] \rightarrow \tilde{y}_{ik}$$

- Identify model components with a causal role in the prediction of y_{iq}

- Mean task-conditioned activation $\bar{a}_{\ell j}^t = \frac{1}{|P_t|} \sum_{p_i^t \in P_t} a_{\ell j}(p_i^t).$

- Causal indirect effect (CIE) to recover y_{iq}

$$\text{CIE}(a_{\ell j} \mid \tilde{p}_i^t) = f(\tilde{p}_i^t \mid a_{\ell j} := \bar{a}_{\ell j}^t)[y_{iq}] - f(\tilde{p}_i^t)[y_{iq}].$$

- Repeat this process for each attention head in f , for all layers ℓ , and all head indices j - highest CIE over all layers

$\rightarrow$ in a set $\rightarrow A \rightarrow \boxed{v_t = \sum_{a_{\ell j} \in A} \bar{a}_{\ell j}^t} \rightarrow \text{Function vector}$

Function vectors → safety amplification

Function vectors - causal mediation analysis
(Todd et al., ICLR 2024)

- Transformer model f
- ICL prompt $p_i^t \in P_t$ dataset representing task t

$$p_i^t = [(x_{i1}, y_{i1}), \dots, (x_{iN}, y_{iN}), x_{iq}] \rightarrow y_{iq}$$

$$\tilde{p}_i^t = [(x_{i1}, \tilde{y}_{i1}), \dots, (x_{iN}, \tilde{y}_{iN}), x_{iq}] \rightarrow \tilde{y}_{ik}$$
- Identify model components with a causal role in the prediction of y_{iq}

- Mean task-conditioned activation $\bar{a}_{\ell j}^t = \frac{1}{|P_t|} \sum_{p_i^t \in P_t} a_{\ell j}(p_i^t)$

- Causal indirect effect (CIE) to recover y_{iq}

$$\text{CIE}(a_{\ell j} | \tilde{p}_i^t) = f(\tilde{p}_i^t | a_{\ell j} := \bar{a}_{\ell j}^t)[y_{iq}] - f(\tilde{p}_i^t)[y_{iq}].$$

- Repeat this process for each attention head in f , for all layers ℓ , and all head indices j - highest CIE over all layers

→ in a set → $A \rightarrow v_t = \sum_{a_{\ell j} \in A} \bar{a}_{\ell j}^t \rightarrow \text{Function vector}$

Safety amplification (SA) in SafeInfer

Prompt $p = \{(q_1, a_1), (q_2, a_2), \dots, (q_n, a_n), q_{n+1}\}$

$(q_i, a_i) \rightarrow$ Harmful question, safe answer pairs

$$\text{attn}'_{lj} = \frac{1}{|P|} \sum_{p \in P} \text{attn}_{lj}(p)$$

$$\text{SV} = \sum_{\text{attn}_{lj} \in A} \text{attn}'_{lj}$$

$$h'_l = h_l + \gamma * \text{SV}$$

M_t

M'_t

(Safety amplified using FV).

Q_1

Model arithmetic

Model arithmetic (DeKoninck et al., ICLR 2024)

	Model Arithmetic	Optimization Problem
Linear Combination	$\sum_i \lambda_i Q_i$	$\sum_i \lambda_i D_{\text{KL}}(P Q_i)$
Union	$\text{union}(Q_1, Q_2)$	$D_{\text{KL}}^{[I_1]}(P Q_1) + D_{\text{KL}}^{[I_2]}(P Q_2)$



$$I_1(x) := [Q_1(x) > Q_2(x)] \text{ and } I_2(x) := 1 - I_1(x)$$

$$D_{\text{KL}}^{[I_1]}(P||Q_1) + D_{\text{KL}}^{[I_2]}(P||Q_2) \rightarrow \sigma(\max(\log Q_1, \log Q_2))$$

$$Q_1 - \lambda \text{union}(Q_1, Q_2)$$

- Resulting distribution only biases tokens for which $Q_2(x) > Q_1(x)$
- Else we recover the original distribution Q_1 (up to a normalization constant).
- Allows us to keep the resulting distribution as close as possible to the original distribution Q_1 , while still biasing away from Q_2

Model arithmetic → controlled generation

Model arithmetic (DeKoninck et al., ICLR 2024)

	Model Arithmetic	Optimization Problem
Linear Combination	$\sum_i \lambda_i Q_i$	$\sum_i \lambda_i D_{\text{KL}}(P Q_i)$
Union	$\text{union}(Q_1, Q_2)$	$D_{\text{KL}}^{[I_1]}(P Q_1) + D_{\text{KL}}^{[I_2]}(P Q_2)$

$\downarrow$
 $I_1(x) := [Q_1(x) > Q_2(x)]$ and $I_2(x) := 1 - I_1(x)$

$$D_{\text{KL}}^{[I_1]}(P || Q_1) + D_{\text{KL}}^{[I_2]}(P || Q_2) \rightarrow \sigma(\max(\log Q_1, \log Q_2))$$

$$Q_1 - \lambda \text{union}(Q_1, Q_2)$$

- Resulting distribution only biases tokens for which $Q_2(x) > Q_1(x)$
- Else we recover the original distribution Q_1 (up to a normalization constant).
- Allows us to keep the resulting distribution as close as possible to the original distribution Q_1 , while still biasing away from Q_2

safety guided decoding strategy (SGDS):

$$D_{\text{KL}}^{[I_1]}(\mathcal{C} || M'_t) + D_{\text{KL}}^{[I_2]}(\mathcal{C} || M_{usf}) \left. \begin{array}{l} \text{where } I_1(x) = [M'_t(x) > M_{usf}(x)] \\ I_2(x) = 1 - I_1(x) \end{array} \right\}$$

$Q_1 - \lambda \text{union}(Q_1, Q_2)$

$$\mathcal{C}(x) = \sigma(\max(\log M'_t(x), \log M_{usf}(x)))$$

$$\begin{aligned} \check{M}_t^{sf} &= \check{M}_t' - \lambda \cdot \sigma(\max(\check{M}_t', M_{usf})) \\ &= \check{M}_t' - \lambda \cdot \mathcal{C} \end{aligned}$$

Results

Metric →
$$\text{ASR} = \frac{\text{\# responses not aligned with Safety}}{\text{\# input queries to LLM}}$$

Baselines

	DangerousQA	AdvBench	HEx-PHI	NicheHazardQA	TechHazardQA	HarmEval
Base model	12.50	20.00	49.09	31.55	43.00	21.63
SafeDecoding	5.00	4.92	6.36	2.77	9.10	6.00
Self-CD	5.50	3.30	4.20	8.79	20	9.45
SA	4.00	14.62	23.64	19.92	45.57	14.55
sGDS	5.50	1.92	5.45	2.34	8.85	1.82
SAFEINFER	3.00	2.69	3.64	1.94	6.14	1.09

ASR of harmful responses for the Llama-2 model across all datasets for the simple prompt setting.

TechHazardQA				
	Instruction-centric		CoT	
	Llama-2	Mistral	Llama-2	Mistral
Base model	86.85	57.57	89.14	41.42
SafeDecoding	27.00	-	19.29	-
Self-CD	40.29	55.43	36.14	40.14
SA	87.71	57.86	88.57	49.28
sGDS	28.28	47.85	16.85	36.28
SAFEINFER	16.57	46.28	14.85	34.85

ASR of harmful responses for instruction centric and instruction-centric CoT prompts.